

DD2425 Robotics and Autonomous Systems Final Report

Christoffer Jansson

Alan Khudur

Dmitrij Lioubartsev

Michal Staniaszek

Yavor Trasiev

December 15, 2014

Abstract

In this report we describe our hardware and software solutions for the house service robot task set in the DD2425 Robotics and Autonomous Systems course. The task required the construction of a robot from a limited set of materials, and the implementation of a software system using the Robot Operating System (ROS)[8] framework. We implemented a control system for motion in the maze, including wall following, a vision system to make use of the Primesense RGB-D camera, and additional systems for mapping.

1 Task Specification

The website for the course specifies the task as follows:

Your robot is the new service robot in someone's house. The new owner has just turned on the robot and given it a few minutes to have a look around in the new environment. Your robot should take this chance to learn as much as possible about the environment so that it can be as good as possible in future tasks. It should learn to find its way and it should detect and remember where certain objects are.

This specification is a description of a real world task that might be performed by a robot. Since we had only two months to implement the system, the actual task that had to be performed was somewhat simplified. The “house” was replaced by a maze made of straight pieces of wood, with all of the walls having either a horizontal or vertical orientation — there are no diagonal walls. The robot should detect and remember the location of various brightly coloured shapes, seen in Figure 5.

The task can be broken into two phases, each of which has a distinct purpose. In the first phase, the robot must explore the environment and learn where objects are. In the second phase, which is not described explicitly in the specification, the robot should return to the previously discovered object locations and “fetch” the objects.

Although these tasks would be trivial for any human, for a robot to do them autonomously is a very demanding task, even in a restricted environment such as the one we will be operating in. To complete the first phase, we must be able to move the robot, which requires the implementation of controllers which allow the robot to move based on demands on angular and

linear velocity. A wall following system is also required, to use the structure of the environment to explore it. This requires the use of sensors to detect walls to the side and in front of the robot to prevent collisions, and detect when it is possible for the robot to turn. A vision system which can detect objects and correctly identify them is also needed. A map must also be constructed and stored so that the position of objects can be remembered for use in the second phase. During the second phase, some sort of path planning is required to move efficiently between the different objects. The ability to localise within the created map is also necessary in order for the robot to know the position of objects and walls relative to its own location. In addition, a way of navigating between specific locations in the map is required.

To complete the full task, we had to design and build a robot, and write software for all of these subsystems. The subsequent sections describe our approach to each problem and how we solved it, including some of the ideas that we did not use, and the reasons for that, as well as some analysis of the performance of each of the subsystems.

2 Hardware

We decided early in the project to build the robot onto a circular platform, so that no part exceeds its radius. This was done in order to have a failsafe mechanism in which the robot could rotate around the z-axis without much problems if it would get stuck in corners and move on. Underneath the circular platform the robot had its two motors with wheel, two caster wheels, six IR sensors and the battery mounted. On top of the platform, the Intel NUC, Arduino board and the Primesense camera were mounted. See Figure 1.

As nice as it is to have a circular body that does not get geometrically stuck, there is a lot of bulk to it, which is definitely felt in the tighter parts of the maze. One could be forgiven to think it is better to solve this problem by code instead of using the physical shape, and that's probably true. However we wanted to implement the solution as quickly and simply as possible. Also, our idea for having symmetry meant having two support caster wheels. The problem with having four support points is having them all aligned perfectly. This warranted adjustable wheel carriers which had to either be a simple L-shape or would have needed intricate holes to facilitate adjustment. We went with the former, which resulted in them constantly bending up and since the caster wheels have a significant offset this

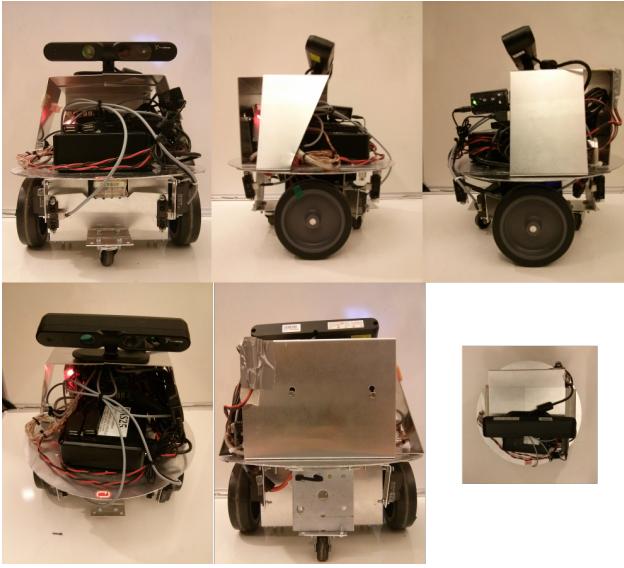


Figure 1: Views of the robot from different angles.

meant we couldn't steer. We fixed this issue by adding an additional bracket.

2.1 IR Sensors

The robot used infrared sensors in order to detect obstructions in the environment. For this particular project, the obstructions were walls in the maze which the robot would traverse in. Six sensors were used - two long range sensors of type Sharp GP2D120 and four short range sensors of type Sharp GP2D12. The long range sensors were specified for a range of 10-80 cm while the short range sensors for a range of 4-30 cm. The sensors were positioned in such way that two short range sensors were mounted on the left respectively on the right side and the remaining two long distance sensors on the front side of the robot. The purpose of this layout was to use the side sensors to detect, follow and avoid collisions with the walls in the maze while the front sensors would be able to detect any obstructions such as object or walls in front of the robot. The front sensors were positioned in an angle towards each other so that any thin obstructions which could lie in between these two sensors could be detected (see top left-most image in Figure 1).

The output of the sensors were numbers ranging from 0 to 550 for the long range sensors and 100-550 for the short range sensors. Each number corresponded to the distance from the sensor where a higher number indicated that an object were closer to the sensor. The correlation between the raw data and the distance was exponential and the raw data contained a lot of noise, making readings of the raw data in real time neigh impossible. In order to solve this problem the sensors had to be calibrated. This was done using a movable wall which was positioned orthogonally to each sensor. For each distance (between 4-30 cm for the short range sensors and 10-80 cm for the long range sensors) data was recorded using rosbag for a couple of seconds. The wall, starting from the closest position each sensor could register, was moved 2 cm backwards in between each reading in order to minimize the amount of data

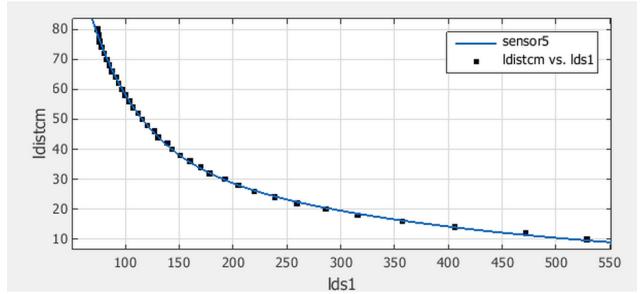


Figure 2: Distance in cm as a function of the raw sensor data given by a long range sensor.

points as it would be too tedious to process later on. When the collection of the data was done, each rosbag was extracted of its data and put in a separate text file. The text files were read to Microsoft Excel in where all data for each distance and sensor was extracted. This extracted data was later read to MathWorks MATLAB where the mean value for all collected data for each distance and sensor was calculated (we discussed if we were to use the mean or the median, but decided on the mean for now since the difference was minimal). For each sensor, each distance point we used could now be plotted as a function of the calculated mean values and fitted with MATLAB's built-in tool cftool using a polynomial curve which later, after a couple of recalibrations, was changed to an exponential curve due to problems of unexpected curvature between several calculated points. All exponential curves which were fitted were of second degree in form of:

$$f(x) = ae^{bx} + ce^{dx} \quad (1)$$

MATLAB was kind enough to calculate and return the values of the constants a , b , c and d . This was later added to a C++ node in which raw output data from each of the sensors were read, added as the x -value in equation 1 and published all in real-time for later use in the wallfollowing and exploration controllers.

Unfortunately, the sensors given to us were of low quality and had a tendency to switch the returned values for every distance every day, forcing to be recalibrated. At first, this issue was battled and several recalibrations were made, but later on this issue was neglected and instead adapted to in the wallfollower. Another problem was that the sensors had problems outputting correct distances close to the limit of the sensor ranges which was solved by putting the sensors closer to the center of the circle (see Section 4.1).

3 Controllers

Controllers are needed to regulate all the motion-related systems not only on the low level (e.g. direct motor control) but on a higher level as well — for example, a wallflower controller that controls the input of the motor controller in order to align the robot to the wall. This is not always the optimal solution, since each controller introduces some lag and imprecision, so the final result may be of poor quality no matter the amount of controller tuning. This is why some of

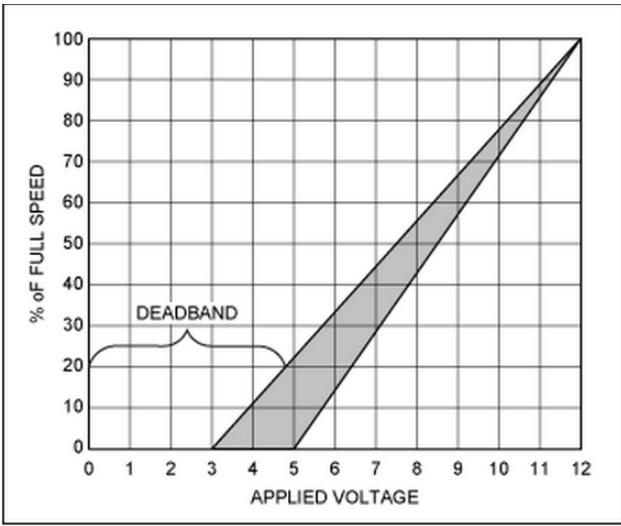


Figure 3: Nonlinearity of DC motors: hysteresis and deadzone.

the controllers used were direct controllers, i.e. they had a direct control of the motors. Unfortunately, this means that there was no way to have a neat single motor controller that runs in the background and awaits commands, but rather a bloated node with two distinct controllers in it — one for roaming and wallfoling and one for precision turning.

3.1 Motor Controller Node

The structure of the node is such, that at any given moment, the robot is either static, or in one of two modes: free-roaming mode or precision turning. It consists of two controllers, each of which controls one of the motors. Obviously, the difference in the velocities of each wheels defines the dynamics of the robot, so this is used through a motion model to either control the rate of turning of the whole robot while keeping a constant linear velocity, or the total distance each wheel travels to achieve a rotation of the robot with zero linear velocity. This is node that accepts a Twist message which defines the linear speed and angular velocity of the robot in free-roam mode, and a turn value, which unless set to 0, defines a precision turning reference in degrees. After each transition of the states the memory of the controller is reinitialized to make sure there is no windup. That said, there was a certain set of cases where there windup occurred frequently which meant we needed to ad anti-windup measures and take that into account during the tuning. The node also communicates with the Master node (Explorer) by conveying a message for whether it has finished turning or not.

3.2 Tuning

Each mode was tuned separately, because of the non-linearity of the motors. As seen in Figure 3, the hysteresis and deadband of the motors, accentuated by the lossy non-linear gearbox makes control of the left and right motors highly dependent on each other. Another complication was the fact that the reference for the roaming mode for each motor was the linear ve-

locity of the wheel, and for the precision turning mode the reference was the distance traveled. All this meant that two set of tuning parameters had to be derived, taking into account the if the direction of the wheels was the same or opposite. A nearly perfect control was achieved by using a bias on the control output which was carefully selected by the characteristic of each motor in each direction of rotation. Although there are hundreds of pre-determined methods for controller tuning, most of them require a highly sophisticated model which could not be justified by our team — it would have taken more than two weeks to complete, while the “by eye” tuning approach which relies on experience and understanding of the control algorithms and motor characteristics took much less. Unfortunately there were more problems:

1. the motors deteriorated to the point where their output speed was up to 30% lower for the same voltage. This process was continuous and we had to tune at least 3 times to compensate for it.
2. The battery position really affected the weight distribution — because it was not completely fixed the weight would shift from the front to the rear caster wheels, causing massive problems with calibration due to uneven caster wheel carriers.

3.3 Wall Follower

The type of controllers we used we mainly PID (more here), however the wall following worked just fine with only a P controller due to the inherent inertness of the system and the fact that it controls the reference input of the main motor controller node. The controller uses the difference between the front and rear sensors on a particular side to control a reference value for the main motor controller.

4 Wall Following and Exploration

4.1 Wall Follower

Figure 4 shows the test maze that we had to work with. The robot was designed to be round to avoid situations where it got stuck in corners or unable to turn in any position. This design although well intended meant that the robot had to be quite wide which leaves little room for error when turning. The robot is 28 centimeters wide at any point and the maze is approximately 40 cm wide meaning that the robot have 5-6 cm on each side when going through a corridor in the maze. This means that the wall following have to be quite accurate and not stray too far from the straight path in the middle. Hitting or bouncing into the walls will throw off the odometry of the robot and result and adversely affect the mapping. The original wallfollower node was written to subscribe to raw data from the arduino sensors for wall following and encoders for determining when to stop turning. The wall following part was done following the left wall with a simple P controller and was simply based on following the wall

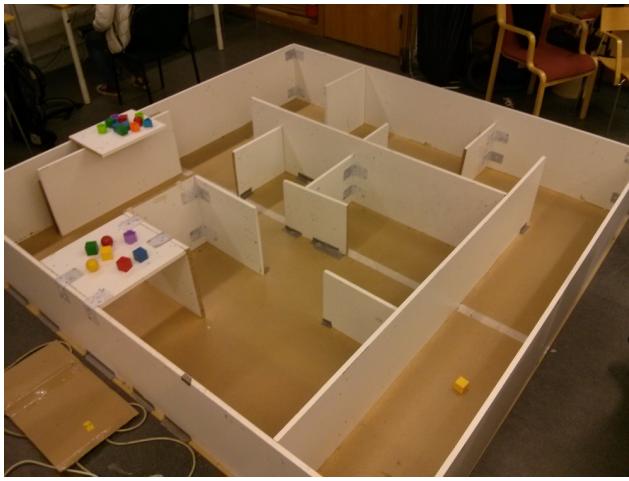


Figure 4: The maze

and when the left wall ends try to turn left. To determine when to turn a state machine was written using a integer. This solution was working poorly so a new version wallfollower2 was written. Wallfollower2 used a byte to store information about which sensor had contact with the wall resulting in a value between 0 and 63. This value was then used to assign a function pointer reference from an array of function pointers to a function pointer that was called each cycle. This had the disadvantage that you had to make 64 functions to get all possible states to work which seemed ineffective so another program was written named simple explorer.

4.2 Explorer

The simple explorer uses the same principle of storing sensor data in a byte but had a separate integer that stored the state. This reduced the number of necessary state functions from 63 to 5. These states are if you dont have an contact on the left side or the right side go forward. If sensor reading comes from the front facing sensors turn 90 degrees to the left or right if there is an obstacle registered by the left facing sensors. If there is contact on both left facing sensors follow the left wall and if thats not the case try following the right wall if the right facing sensors have contact. This however meant that it was no longer a real left side wall follower so a complex explorer was written to solve that situation. The complex explorer adds three states that make sure that the robot always follows the wall on the left side. The first state is if there is nothing on the left side sensors turn 90 degrees to the left. After that using the left facing sensors to scan for a wall to follow or if an edge of a wall is registered turns 90 degrees left and then drives forward. The final state that was added was to turn 180 degrees if walls are registered by all front sensors. This should result in a robot capable of traveling through any maze however there were issues. The most severe issue was false sensor readings which was solved by gathering a couple of sensor bytes after a change has been made before changing state. Another limitation was the placing of the front forward facing sensors which was placed 7.5 centimeters from the roll axis of the robot. This means that the robot can only

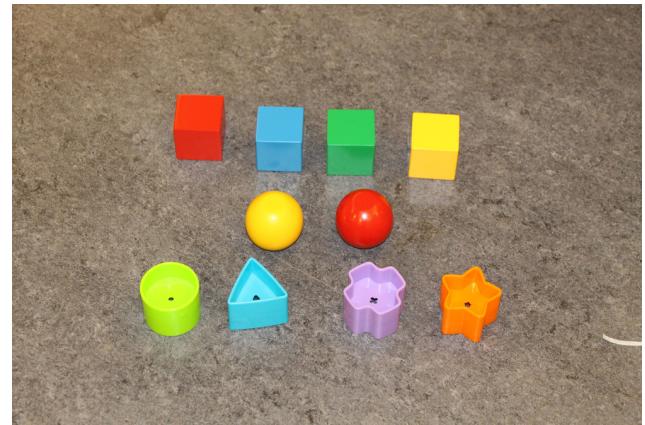


Figure 5: Objects to be detected within the maze.

discover walls in front of it if it was within a 15 centimeter from around the roll axis with a 28 centimeter wide robot meaning that the robot had a blind spot of 4 centimeters on either side facing forward. This was solved by stopping quite far from the wall for turning making sure that the robot could follow the entire test maze to the end in one go.

5 Vision

5.1 Observations About Objects

The object set can be seen in Figure 5. Each object only has one color, and no texture features. There are objects of the same color but different shapes and objects of the same shape but different colors. Because of the lack of texture, feature-based detection[3] will not work very well. However, because all objects are all mono-colored, some sort of detection based on color is appropriate. Having only color detection is not enough though, as objects with the same color need to be differentiated through shape. This can be done with either looking at contours or utilizing the depth part of the Primesense camera and looking at the object shape. We chose the shape method.

5.2 Color Detection

The objects can be separated into 6 different colors: blue, green, red, yellow, orange and purple. The green objects and the blue objects are not exactly the same color but they are similar enough to be grouped together.

For the color detection the RGB sensor of the Prime-sense camera was used. RGB sensor is a buzzword and it is more commonly known as simply camera. The camera outputs an image stream. For each frame, we get an image encoded as an RGB image. RGB stands for Red Green Blue and is the most common color representation model. An image consists of many pixels. Each pixel uses three values (usually between 0 and 255) to represent the amount of each color. For example, (0,0,0) is black and (255,255,255) is white. Each combination of RGB values yields a unique color, so this representation allows for 256³ different colors. However, obviously colors with very similar values look

very similar. Changing just a few points in each value will barely be detectable by the human eye. So while we can say that each object consists of only one color, that is technically not true. Things like dirt, shadows and background/environment light will change the objects pixels values.

Because of this, we need to define a range of pixel values which represent each color. In other words, each color will have a model. Then during detection, each pixel will be compared to the model to get some number on how similar a pixel is to the color. This is done for every pixel in the image in each frame. Then, using a threshold, a threshold image is created. Basically, every pixel similar enough is set to 1, and otherwise 0. Then we check if there is a big blob somewhere, and if so, we say that we found an object of this certain color. We used our own algorithms combined with OpenCV algorithms to build the color detection system.

5.3 Model Training

The simplest color model would be to just put some threshold values directly on the RGB values of each pixel. That has certain issues, as it is hard to know exactly where the different thresholds somewhere, and if so, we say that we found an object of this certain color. We used our own algorithms combined with OpenCV algorithms to build the color detection system.

Model training The simplest color model would be to just put some threshold values directly on the RGB values of each pixel. That has certain issues, as it is hard to know exactly where the different thresholds should be. A thing to note about the RGB encoding is that it contains data about light intensity. If the ratio between the three values stays the same, but the values change, the color stays the same, but just changes illumination (gets lighter/darker). To make the model less affected by illumination, instead of using RGB representation we transform all pixels to the *rg*-chromaticity representation [7]. In short, given a pixels *R*, *G* and *B* values, we calculate *r* and *g* as

$$r = \frac{R}{R + G + B} \quad (2)$$

$$g = \frac{G}{R + G + B} \quad (3)$$

and use that as the representation in each pixel. This also simplifies the model as it is only 2-dimensional now.

We also tried the representation used in [1]. It did however not yield better results, and was just generally slower, so we stuck to just using rg-chromaticity.

For the model, we used a two-dimensional multivariate Gaussian to represent each color. A two-dimensional Gaussian has in total 6 parameters: the means of each colours (a 2-dimensional vector μ) and the 2×2 covariance matrix Σ which represents standard deviation and co-variation. Using training data, we calculated these parameters for each color using the formulas given in [9]. To summarize: μ is the mean vector, x is the vector of all the training pixels, and N is the number of training pixels. One could say that x is technically a $2 \times N$ matrix.




Figure 6: Training image, object selected (with 1 click), and then pasted onto a black background.

$$\mu_i = \frac{1}{N} \sum_{j=1}^N x_{ij} \quad (4)$$

$$\Sigma_{jk} = \frac{1}{N-1} \sum_{i=1}^N (x_{ij} - \mu_j)(x_{ik} - \mu_k) \quad (5)$$

For the training data, we wrote a quick program with which we could take images with the Primesense camera. At first we used our cell phone cameras but the model was more accurate when we used the Prime-sense. The images were taken in the test maze to get the training data to be as similar to the actual data as possible. The objects were then extracted from the images and put on black backgrounds using Photoshop. Photoshop's quick selection tool made it very easy to select the objects, and it took around 30 seconds per training image. The process is shown in Figure 6. It took more time to save the new images than actually select and extract the objects. Before using Photoshop, we tried doing the extraction automatically by taking pictures of the objects on a white background and using thresholding to extract the object from the image. The extraction was quite fast, but often resulted in parts of the background being included in the resulting image, which was negatively affecting the training data. We decided that accuracy was more important than speed in this case, and instead did the extraction manually. We used approximately 40 training images for each color. We wrote a program which calculated the μ and Σ from all images in a given directory. It assumed that every non-black pixel was from the object, and then just used the above formulas for the calculations. Because this is ran once, before we even run the robot, this does not need to be efficient. This way we had a fairly good framework for getting training data, and we could easily add new training images or remove existing images.

5.4 Detection

When the robot is moving around in the maze, we get an image for each frame. We first convert it to rg-chromaticity representation, and then using OpenCV blur it a bit to lower the impact of noise. Then we calculate the similarity of each pixel to the model (this is done once for each color). For this we simply insert

the pixel values into the Multivariate Gaussian formula

$$pixel = \frac{1}{\sqrt{(2\pi)^2 |\Sigma|}} \exp \left(-\frac{1}{2}(x - \mu)^T(x - \mu) \right) \quad (6)$$

$$= \frac{1}{\sqrt{(2\pi)^2 |\Sigma|}} \exp (-0.5(r_n^2 \Sigma_{11}^{-1} + 2r_n g_n \Sigma_{12}^{-1} + g_n^2 \Sigma_{22}^{-1})) \quad (7)$$

Where $r_n = (r - \mu_r)$ and $r_g = (r - \mu_g)$. $|\Sigma|$ is the determinant of Σ , and Σ^{-1} is the inverse of Σ . Also note that $\Sigma_{12} = \Sigma_{21}$ is always true. This is not very efficient calculations to make for every pixel. Some simple optimizations can be made by moving -0.5 into the big parenthesis. Because Σ and μ never change between runs (they are from the model), things like Σ^{-1} and $|\Sigma|$ can be calculated only once upon start. In fact, the whole big thing left of the exp can be pre-calculated. To make it even faster we calculate the natural logarithm of the expression as well. Then we get

$$pixel = c \cdot (r_n r_n p_1 + r_n g_n p_2 + g_n g_n p_3) \quad (8)$$

Where the p_i is some pre-calculated variable. This means, that for this part we only use 7 multiplications and 2 additions.

The next step in the algorithm is to make the threshold image. The threshold was set without any training by trial-and-error. Because we calculated the logarithm of the pixel value, we have to set the threshold to the logarithm of the threshold. This is done once during startup so it is ok. The issue with doing the threshold method is that we lose the similarity data. The strong part with the gaussian model we use is that you get some sort of probability for each pixel being that certain color. This information is lost with the thresholding. We experimented with some ways of utilizing this data, but never got anything successful working. With the thresholding, we could utilize the OpenCV algorithm findContours.

findContours does what the name implies. It finds the contours in a binary image. We then take a simple approach and simply check the size of the biggest contour. If it is larger than a certain threshold, we say that the image contains this particular color.

Note that everything above from blurring, is done once for each color every frame. So we can detect multiple different colours in the same image.

5.5 Shape Detection

We made a key observation about the object dataset: objects of the same color are either a cube or non-cube. This means, that if we could distinguish between a cube and non-cube we could distinguish between all objects. So we needed the shape detection to look for cubes. Another observation we made was that the cube is the only shape which has a plane which is parallel to the floor. The circle has one but it is significantly smaller. This crafted the ridiculously simple idea, that we could just look for planes parallel to the floor in the pointcloud. A pointcloud is the one of the outputs from the Primesense. It gives a set of points with x, y, z coordinates, essentially giving us a 3D image. Each point




Figure 7: Purple color detection. Bottom left is the full RGB image. The bottom right is the image with the pixel values calculated from the model (brighter means more likely to be purple). Top right is the thresholded image. Top left is the console writing each frame if the color is detected or not.

also contained color information but we did not use it. We used the library PCL [6] to work with pointclouds.

The first step of the algorithm is to extract the floor. This is done with the RANSAC algorithm which is provided by PCL. It extracts the dominant plane in a pointcloud, which often is the floor, however when the robot drives close to a wall it becomes the wall. It also gives the planes coefficients (on the form of $ax + by + cz + d = 0$). We identified the typical coefficients when the dominant plane is a floor, and put in a check, so that we do not try and find cubes when the dominant plane was a wall. We also tried a version where we extracted multiple planes, but that did not work better.

Using the extracted planes coefficients and the pointcloud with the floor removed, we ran the PCL algorithm NormalEstimation. This looks at each point, takes its closest neighbors, assumes they are a plane, and calculates that planes normal. This was quite slow on a large pointcloud so we downsampled the pointcloud first (another PCL algorithm). In practice, this meant that all pixels on the top side of a cube (except those close to the edge) would have normals pointing straight up - same as the floors normal! Due to the environment being so restricted, cubes are the only shapes that would get normals pointing straight up, so we could just count the amount of normal vectors pointing up, and if they were many enough we said there was a cube in the image.

To check if a vector pointed upwards, we used the dot product between the floor normal and the vector to get the angle between them. If the angle was low enough (say, 5 degrees), the vector is said to be parallel with the floor.

5.6 Vision Master

Using the data from the color and the shape detection, we could identify which object was in the image, if there was one. The color was used primarily to find the object, and then the cube data (just a boolean) would be used to distinguish which object it was of that color (if needed). This was done in the vision master part of

the program. If a certain object was seen for a certain amount of frames in a short time, that object counts as found. Each object can only be found once, as the vision did not keep track of where a certain object was found.

To tell the mapping where the object was found, we used the depth data from the Primesense, the Prime-sense Field of View values, the contour centers from the color detection and the height the camera was positioned at (approx. 27 cm), we could calculate the offsets in different axes where exactly the object was positioned relative to the robots position.

This detection works fine (in theory!) if there is only one object in view. This will not always work properly with more than one object in view. To fix that was not our highest priority. But we did try a variation where we utilized the position of where the object was found, and only looked for a cube in that particular area. This could then be done for each color found, essentially meaning that we could detect multiple different objects in the same view. Since the color detection only took out the biggest contour, we still could not detect both if there were two objects of the same colour in the same view. That is such a rare case so we never bothered fixing it.

5.7 Did it work?

Well, kind of. Sometimes it worked great and other times not.

5.7.1 Color

The color worked really well for purple, blue and green. Especially purple. I would be surprised if any other group was as good as us on the purple cross. See Figure 8. From over 1 meter distance, while the robot was moving, it sees the purple cross while it is very badly illuminated. Blue and green are not as insane but are quite good.

The color yellow was not detected very well. We never adjusted the parameters very well. But we were afraid to play with them too much as the floor was of a similar color. Out in bright spots, yellow was detected decently.

Red was always detected fairly well. The problem was that red and orange were extremely similar. We did not expect that. Red objects would often be misclassified as the orange star, and vice versa. By the time we found this out it was too late to do anything about it. Well, we tried taking training pictures more carefully, but that made the models even more similar.

5.7.2 Shape

We did not spend nearly as much time on shape detection as on color detection, and it shows. In theory, even though the idea is very simple, it should work. And it did work in our initial tests. We could actually detect all 10 objects. The thing is, that the robot was standing still. And the objects were on the middle of an open floor.

When driving in the maze, there is more noise, and the objects typically stood next to walls. Especially

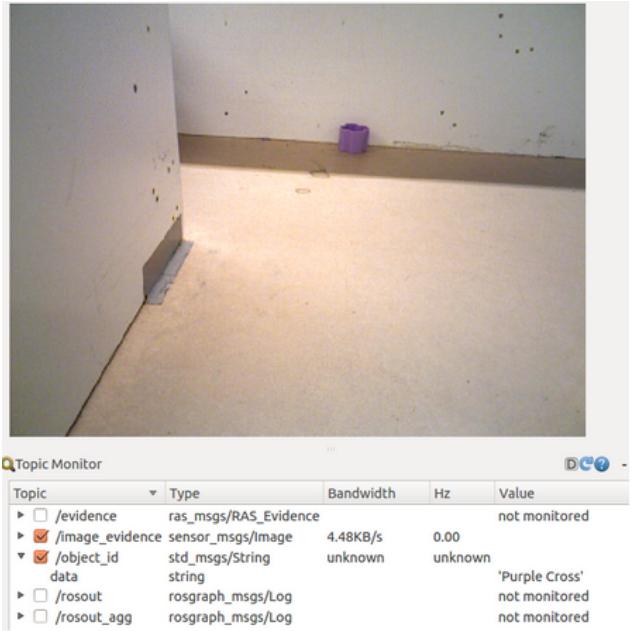


Figure 8: Purple object detected, from the evidence viewer.

them standing close to walls really messed up the cube detection. In attempt to fix this, we tried things like removing several dominant planes (already mentioned) and looking for parallel vectors only in the approximate area where the color detection found the object (also mentioned). To try and reduce noise was actually the main motivation to try those things. Unfortunately nothing of this worked, and then we were out of time.

5.7.3 Contest

Well, we never bothered adjusting the color detection to the vastly different light conditions in the Ljusgården compared to the lab. Because of that, nothing in the detection worked properly and we misclassified a lot of objects. Of course, the purple cross was still detected correctly (when we tested in the break).

6 Mapping

In order to perform the task required in the second phase more efficiently, we need a map in which the location and type of object is stored, as well as the locations of walls and other potential obstacles. To make the map we use only the IR sensors. Using the Primesense as well might have given us more data to use, but we didn't have time to implement anything to do this.

6.1 RANSAC Method

Since the map has a regular structure, we decided to try an approach which exploits the regularities and minimises the effect of sensor noise and odometry drift on the resulting map. We thought that most of the issues from odometry would come from the turning. Instead of recording constantly, we only record data when the robot is moving along straight line segments. During turning, the recording is stopped, and the odometry

zeroed. In theory, this prevents the build-up of error, and should mean that the individual segments are closer to the actual position of the map that was traversed. Since each individual segment is independent, some post-processing needs to be done in order to connect the segments to each other. This is done by storing the turn direction at the point where one segment ends and the other begins. The total length of the segment is also extracted from the odometry. Each measurement is paired with odometry readings at the same point. Detected objects are also stored with an offset from the robot frame.

6.1.1 Line Extraction

To exploit the regularities, we use the built in PCL RANSAC [4] on the points generated by computing the coordinates of each of the four sensor readings and then translating them by the corresponding odometry reading. The set of inliers that lie on the line model received from RANSAC is removed from the original point cloud, and the process is repeated on the remaining points. Once this is done, we receive a set of lines which represent the walls that were detected in the segment. Doing this for every segment, we have the set of lines which should give us all the detected walls in the map. However, the lines must still be translated and rotated into the correct position relative to the start position — each segment is saved with its start at the origin, and with no rotation. In addition, RANSAC gives line equations, and not line segments, so the line segment must be extracted from the point cloud. To do this, we look through the inliers point cloud received from RANSAC to find the points with maximum and minimum x, y coordinate values. These must be the start and end points of the line. Once the line segments are extracted, we use the stored turn and odometry information to rotate and translate the end points of the line segment into the correct position.

We did not know whether this would be a computationally intensive process, and so instead of processing the data in real-time, it is stored to be processed after the end of the first phase. The map is not strictly necessary in the first phase, since we are only required to explore and discover objects. Of course, having a map is beneficial when new areas need to be explored, but we decided that this was not a problem.

6.1.2 Map Construction

To construct the map, we make use of the regularity in the map. We assume that adjacent walls are orthogonal to each other, and that they all lie along the x or y axes. Using the lines from RANSAC, we perform some post processing to align them to the axes. Lines are aligned by taking the average of the closest coordinates. That is, if the distance between the x value of the start and end points is smaller than that of y , we take the average of the x values and set the aligned line x to be that average. This makes the line parallel to the y axis. The map is an occupancy grid map, where each cell can be either unknown, occupied or unoccupied. We fill cells on a segment-by-segment basis. First, we find the bounds of the lines which make up

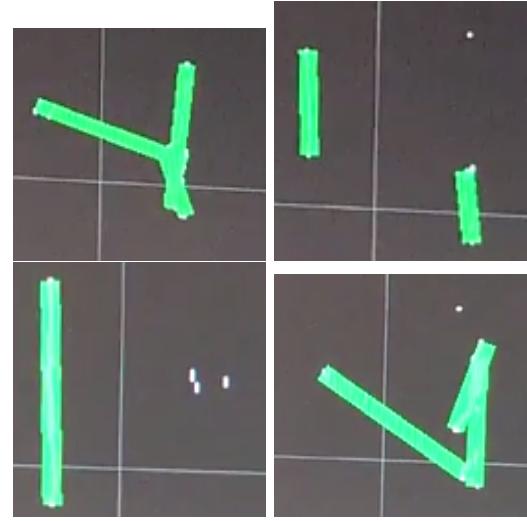


Figure 9: Lines extracted from the IR sensor data using the RANSAC method, from a single segment.

the segment, which has already been translated to its correct location. We then make a bounding box polygon out of these bounds, and stamp it onto the map, setting all the cells in that region to be unoccupied. Then, we do the same for each line in the segment, setting the cells to occupied.

6.1.3 Problems

Examples of the lines extracted by the initial RANSAC process can be seen in Figure 9. Examples of the result of the rotation and translation process based on turn and odometry data can be seen in Figure 11.

There were numerous problems with the RANSAC approach. Firstly, there were some issues with receiving odd sensor readings. This was mitigated by filtering the sensor values with a minimum and maximum threshold. Some valid readings still caused problems with the line extraction — these points would not be removed, and sometimes resulted in lines which were much longer than they should have been, because the point was on the line model. There was also the problem of RANSAC extracting lines with very few points remaining in the point cloud. This issue was solved by adding a threshold on the minimum proportion of points in the point cloud required for the algorithm to be run. Even with these additions, some problems remained. Extraction would find diagonal lines, like in Figure 9, and outliers would cause lines to be extended as before. Sensor readings for maximum distance were being ignored, which meant that regions of free space were not being properly defined when lines were used to construct the map. Some of these problems could be mitigated by outlier removal, splitting data into the left and right-hand data and running extraction on those separately, and some post-processing on the extracted lines, such as checking gradients. However, we were unable to implement any of these ideas before the milestone deadline. Compounded with the issues with map construction, the RANSAC method was abandoned and we moved on to trying to construct a simpler topological map.

The map construction was also too simplistic. The

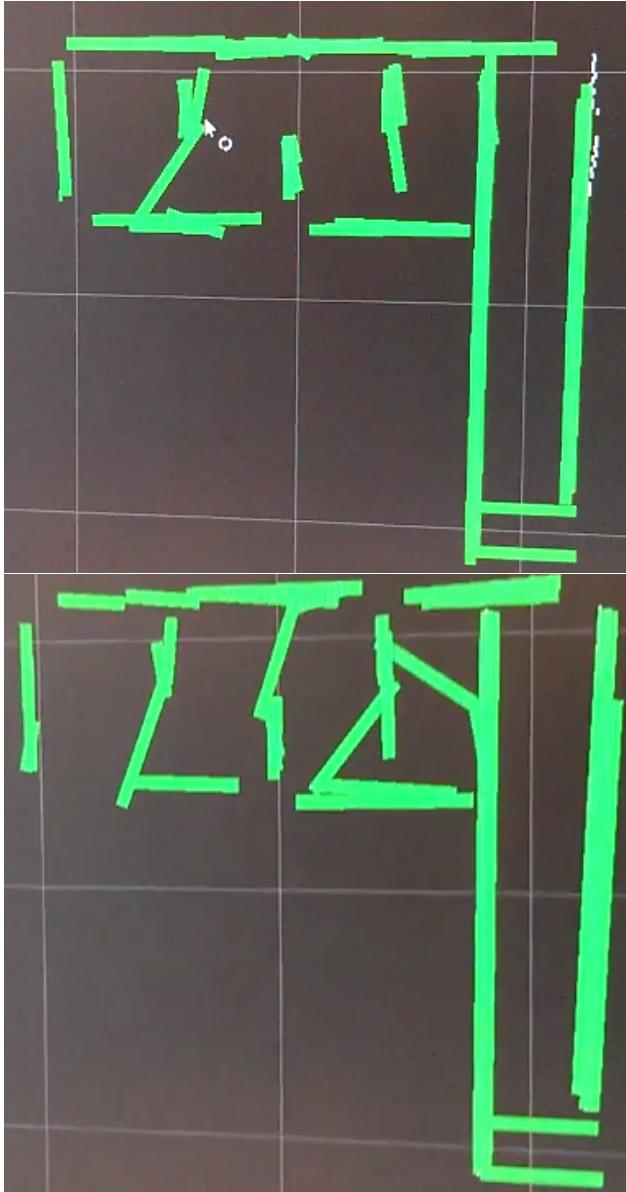


Figure 10: Lines extracted from the IR sensor data using the RANSAC method, rotated and translated according to the turning information and odometry data.

bounding box of the lines is not necessarily the correct region to define as empty space, especially since there may be multiple different “levels” of walls in each segment. For example, in segments where there are three parallel walls, one shorter than the others, and in between them, then this wall will be entirely ignored, and the space it occupies will be assumed to be free space.

The main reason for the unsuccessful implementation was a lack of good task distribution and proper discussion about the approach. If we had discussed the potential problems early on in the implementation, perhaps we would have not chosen to take this approach and waste a lot of time.

6.2 Topological Map

In place of the RANSAC based map, we decided to implement a basic topological map which would place nodes at coordinates where the robot made a turn or detected an object. Node coordinates were again based

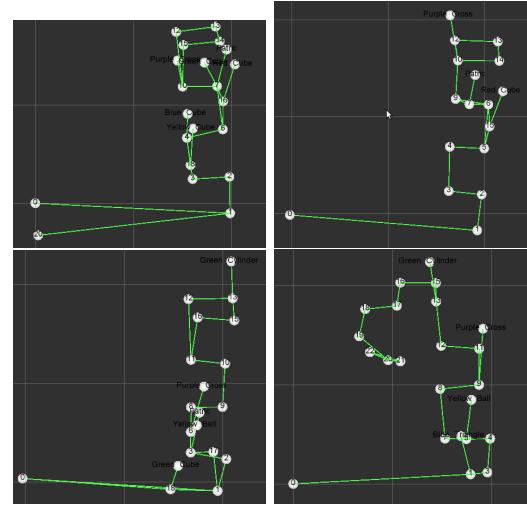


Figure 11: Examples of topological maps constructed from multiple runs in the same environment.

on the odometry. The implementation saves the nodes to a bag file, which can then be read into the system in the second phase. The map can then be used to navigate, and allows for the addition of nodes into the existing map if new areas are explored.

We intended to add functionality to store information about potential areas for exploration, and also to place nodes on the intersection points of lines to improve the efficiency of navigation, but due to time restrictions we were unable to complete these extensions.

7 Odometry, Localisation and Navigation

7.1 Odometry

The odometry was only facilitated through the use of motor encoders and careful controller tuning as to not let any slippage occur. Using the IMU was prematurely decided against since we had problems with the script that was running it. The odometry returned the total amount of distance travelled, the current heading and the absolute X and Y coordinates, all derived using the motion model of the robot, found here.

7.2 Localisation

Localisation was implicit in Navigation as it did not include any filtering, i.e. localization was simply the task of referring the output of the odometry to the beginning of the coordinate system and thus to the node map, since it has metric data. That was done with a simple euclidean distance measure which was used as a threshold of whether the robot is on a certain node or not.

7.3 Navigation

Navigation was implemented as a dual state system - either following the trajectory, described the map, or wallowing for when there were straight walls to help

with the reduction of odometry error. While the wall-following was already well defined, we needed to implement the navigation capability separately. We used a BFS algorithm to find a path between two nodes. After that path following was done by using the difference of the vectors of the current position and that of the target (next accessible node of the path) and having this vector serve as the reference heading of the robot. Thus, there is a controller which makes sure the difference in angle between the heading of the robot and the reference vector is zero. That however added a serious problem - due to odometry errors the path was not always parallel to the walls, which meant that we needed to use a controller which aligns the map to the closest straight wall while the robot was following said map. This ensured that the map will always be aligned to the last seen wall, thereby reducing the error in heading significantly. Please note, that unfortunately we did not have time to implement an explorer which is capable of venturing away from the walls on its own accord, which meant that there was no point in investing time in tuning the tricky cascade of 3 controllers and the node got abandoned. Not all of it, though. A part of it got transplanted into the Advanced Explorer to create a new node under a new name - Fetch. This node uses localization to determine if the robot has reached the position from which a desired object can be seen. Of course it relies on the fact that the map is consecutive and deterministic, since we do not venture away from the walls very often. This meant that the task of fetching could be completed for about half of the objects. The Fetch node was proven to work in the lab, but was not used in the competition due to failure to get a map in the first run.

8 Work Environment

8.1 Version Control

For source control, we used `git`, with repositories on GitHub. Each of us had a separate account on GitHub, but we all committed to the same set of repositories which were created in an Organisation on the site. We had separate repositories for each major subsystem (mapping, controllers, vision, etc.). Everyone committed to the same repositories, without using forks. While using pull requests can have some benefits in terms of allowing team members to look at what has been changed in a specific pull request and decide if changes are required, in practice this sort of review process could only be done with a very focused team that already have quite a lot of experience with `git`. It would also be time consuming to have to do this, and time is already tight enough as it is in the course. We tried to use a version of the branching model described in [2]. The idea is to use a branch `develop` in each repository to develop code, branching off of that to develop new features, in order not to force other members to pull code that is in the process of being developed. The `master` branch is pushed to whenever a milestone (release) is reached. While this may be a good branching model, in practice it is perhaps not so useful if there is only one person working in a sin-

gle repository, as we often had. The good thing about it is that merging `develop` into `master` only when a milestone is reached means that the milestone code is easy to find, and it is possible to make sure that only working code is merged.

There are some advantages to a multi-repository setup — it is possible to work on one repository completely independent of the others, which minimises the number of merge commits and other issues that can be caused by having everything in a single repository. The disadvantage of this is that all repositories have to be pulled independently of each other, and if one forgets to do this it can cause problems due to not having up to date code. The advantages of this setup probably outweigh the disadvantages, however, because a single repository setup would require more use of branches, which have issues of their own.

8.2 System Setup

At the beginning of the project we decided to make separate users on the NUC, each with a separate `catkin` workspace. This worked well for us, as we were able to work on different versions of code without affecting other members' work. This came with the additional benefit that each member could launch their own part of the code for a specific task independently. It did come with the problem that each member had to pull updates to repositories, rather than having a single repository which was up to date all the time. Each user belonged to the same `robo` group, which allowed everyone to access each others' files. A very useful script was one that allowed for execution of `catkin make` from any directory, which saved having to switch between terminal windows or directories. The system setup took some time, but the small amount of time spent at the start was outweighed by the benefits.

8.3 ROS Specifics

We tried to keep to a specific structure for all nodes for consistency. Each node created is an object (although there are no duplicated nodes). This keeps the parameters of the node in its own scope, preventing possible issues with clashes. Each node loads parameters and sets up subscribers in its constructor. After all parameters are loaded, the node calls a function which executes the main loop of the node. This structure makes it easy to see exactly what the node is using by just looking at the constructor. In addition to this, we set up parameter files for nodes which could use different parameter settings. Since modifying extra parameters does not require recompilation, this saves time when looking for the best parameter settings. A global parameter file was used to store information about sensor positions and robot measurements. We make extensive use of launch files as a result of this setup. Each launch file loads the parameters that it requires. As such, it is possible to create a single launch file at the top level which runs multiple nodes. Another parameter file is used to define all topics that exist on the system. This means that one change can change the name of the topic in all nodes which make use of it. This means some ad-

ditional parameter reading in nodes, but it lowers the likelihood of bugs where nodes are not subscribing or publishing to the correct topics.

We implemented a utility library which was a wrapper around the parameter server which would load a parameter into a given variable, depending on the type of the variable. If the parameter was not present in the parameter server, an error would be thrown with information about the missing parameter. This helped to make parameter reading quicker and less prone to errors.

9 Conclusion

References

- [1] Christian Balkenius and Birger Johansson. “Finding colored objects in a scene”. In: *LUCS Minor* 12 (2007).
- [2] Vincent Driessen. *A successful Git branching model*. Jan. 2014. URL: <http://nvie.com/posts/a-successful-git-branching-model/> (visited on 12/15/2014).
- [3] *Feature Detection (computer vision)*. URL: [http://en.wikipedia.org/wiki/Feature_detection_\(computer_vision\)](http://en.wikipedia.org/wiki/Feature_detection_(computer_vision)) (visited on 12/15/2014).
- [4] Martin A Fischler and Robert C Bolles. “Random sample consensus: a paradigm for model fitting with applications to image analysis and automated cartography”. In: *Communications of the ACM* 24.6 (1981), pp. 381–395.
- [5] *OpenCV Official Website*. 2014. URL: <http://wwwopencv.org/> (visited on 12/15/2014).
- [6] *PointCloud Library Official Website*. 2014. URL: <http://www.pointclouds.org/> (visited on 12/15/2014).
- [7] *rg Chromaticity*. URL: http://en.wikipedia.org/wiki/Rg_chromaticity (visited on 12/15/2014).
- [8] *Robot Operating System Official Website*. 2014. URL: <http://www.ros.org/> (visited on 12/13/2014).
- [9] *Sample Mean and Sample Covariance*. URL: http://en.wikipedia.org/wiki/Sample_mean_and_sample_covariance (visited on 12/15/2014).